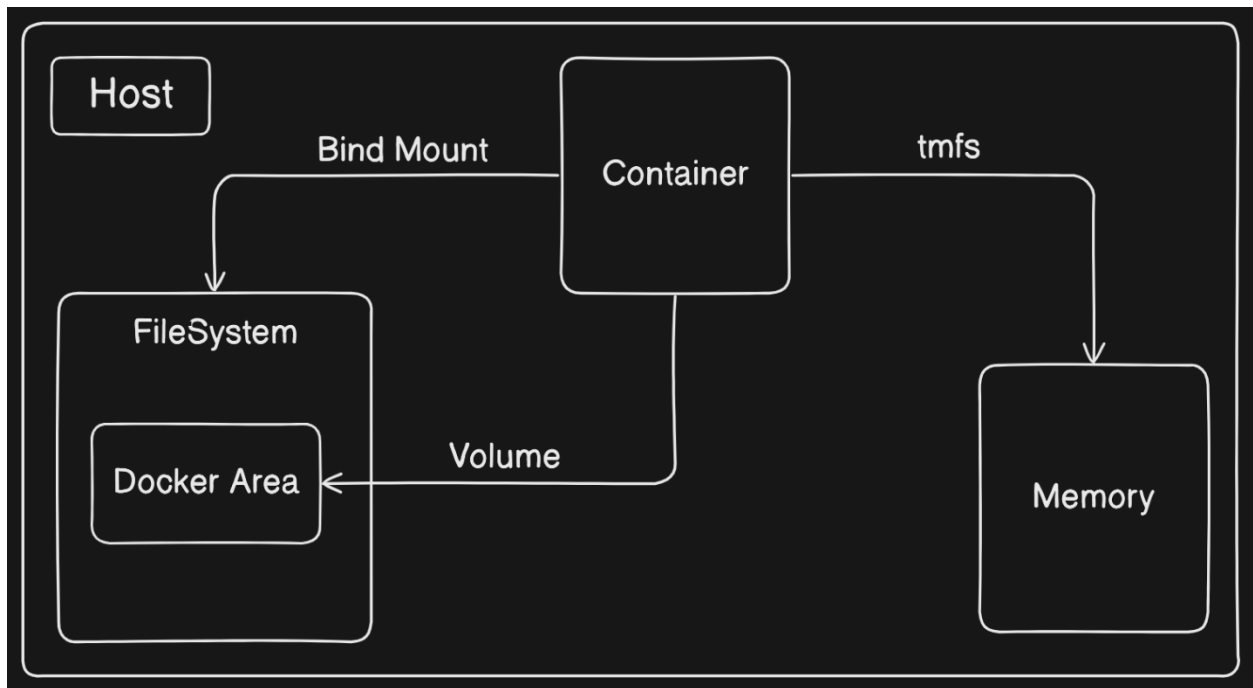


## Working with Docker Volumes

Docker Volume is a mechanism in Docker that allows you to persist data even after a container is deleted or recreated. It provides a way to store data outside of the container's filesystem, so that data is not lost when the container is restarted or deleted.



## Why do we need Docker Volumes?

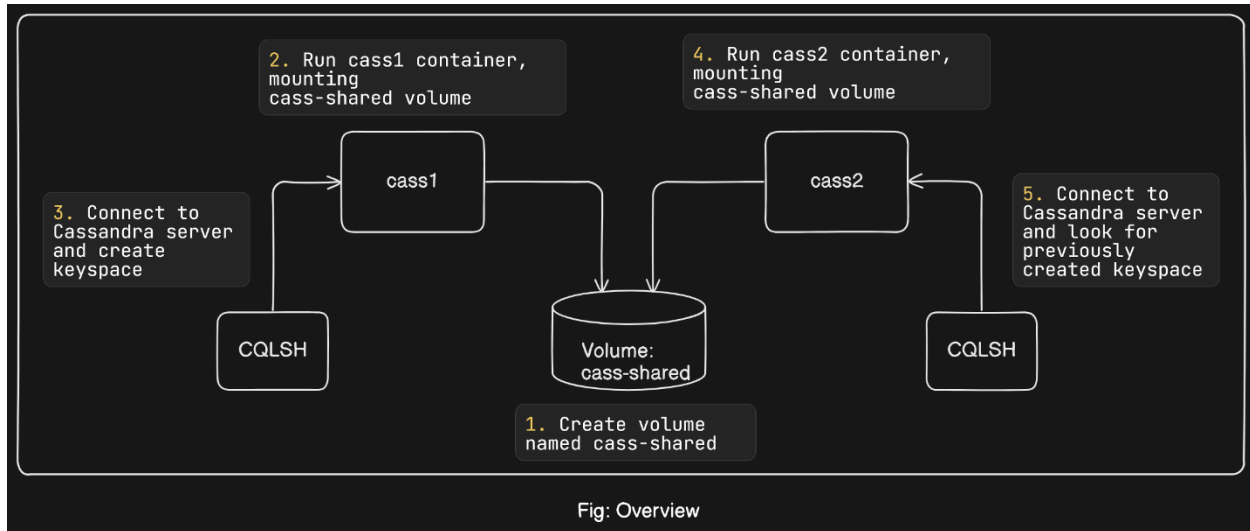
When you run a container, it has its own filesystem, which is ephemeral. This means that any data written to the container's filesystem will be lost when the container is deleted or restarted. Docker Volumes provide a way to decouple the data from the container's filesystem, so that data can be preserved even when the container is recreated.

## How do Docker Volumes work?

A Docker Volume is a directory that is **shared** between the host machine and a container. When you create a volume, Docker creates a directory on the host machine, and mounts it to a directory inside the container. This allows data to be written to the volume, which is persisted even when the container is deleted or recreated.

## Example Scenerio: Using Docker volumes with a NoSQL Database (Apache Cassandra)

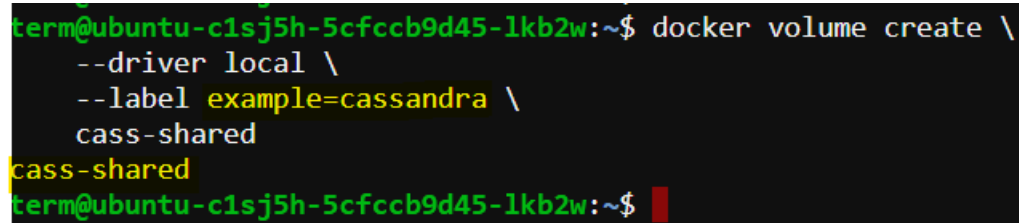
In this scenario, we will use Docker to create and manage a single-node Cassandra cluster. We'll create a keyspace, delete the container, and then recover the keyspace on a new node in another container using Docker volumes. Follow the detailed steps below:



## Step 1: Create Docker Volume

First, we will create a Docker volume that will store the Cassandra database files. This volume will use disk space on the local machine.

```
docker volume create \  
  --driver local \  
  --label example=cassandra \  
  cass-shared
```

A terminal window screenshot with a black background and green text. The prompt is 'term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~\$'. The command entered is 'docker volume create \ --driver local \ --label example=cassandra \ cass-shared'. The output is 'cass-shared'. The prompt returns to 'term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~\$' with a red cursor.

```
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker volume create \  
  --driver local \  
  --label example=cassandra \  
  cass-shared  
cass-shared  
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$
```

### Explanation:

- **docker volume create:** Command to create a Docker volume.
- **--driver local:** Specifies that the volume should use the local driver.
- **--label example=cassandra:** Adds a label to the volume for easier management and identification.
- **cass-shared:** The name of the volume.

## Step 2: Run a Cassandra Container

Next, run a Cassandra container and mount the previously created volume to the container.

```
docker run -d \
  --volume cass-shared:/var/lib/cassandra/data \
  --name cass1 \
  cassandra:2.2
```

```
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker run -d \
  --volume cass-shared:/var/lib/cassandra/data \
  --name cass1 \
  cassandra:2.2
Unable to find image 'cassandra:2.2' locally
2.2: Pulling from library/cassandra
35807b77a593: Pull complete
93d71b8f96bb: Pull complete
b6a858f02311: Pull complete
baff0a18b53b: Pull complete
1abb4abff8c7: Pull complete
ef4d0c404e96: Pull complete
f87427648c5b: Pull complete
b7dfd75caf76: Pull complete
ef996cf87903: Pull complete
Digest: sha256:99f8579937aeea75d0e4f2ad055dfed587f03eebdebfd3666420c1162206503
Status: Downloaded newer image for cassandra:2.2
016b2ce03efd1dab628c47488641f945a2510f36ce4bdfc8c3f7326223957ca9
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$
```

### Explanation:

- **ocker run -d:** Runs the container in detached mode.
- **--volume cass-shared:/var/lib/cassandra/data:** Mounts the cass-shared volume to /var/lib/cassandra/data inside the container.
- **--name cass1:** Names the container cass1.
- **cassandra:2.2:** Uses the Cassandra image version 2.2 from Docker Hub.

### Step 3: Connect to the Cassandra Container

Use the Cassandra client tool (CQLSH) to connect to the running Cassandra server.

```
docker run -it --rm \
  --link cass1:cass \
  cassandra:2.2 \
  cqlsh cass
```

#### Explanation:

- **docker run -it --rm:** Runs the container interactively and removes it after exit.
- **--link cass1:cass:** Links the client container to the cass1 container.
- **cassandra:2.2:** Uses the Cassandra image version 2.2.
- **cqlsh cass:** Runs the CQLSH command line tool to connect to the Cassandra server.

Now you can inspect or modify your Cassandra database from the CQLSH command line. First, look for a keyspace named `docker_hello_world`:

```
select *
from system.schema_keyspaces
where keyspace_name = 'docker_hello_world';
```

## Explanation:

- **select \* from system.schema\_keyspaces where keyspace\_name = 'docker\_hello\_world';** Queries the system schema for the docker\_hello\_world keyspace.

```
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker run -it --rm \
  --link cass1:cass \
  cassandra:2.2 \
  cqlsh cass
Connected to Test Cluster at cass:9042.
[cqlsh 5.0.1 | Cassandra 2.2.19 | CQL spec 3.3.1 | Native protocol v4]
Use HELP for help.
cqlsh> select *
... from system.schema_keyspaces
... where keyspace_name = 'docker_hello_world';

keyspace_name | durable_writes | strategy_class | strategy_options
-----+-----+-----+-----
(0 rows)
cqlsh>
```

Cassandra should return an *empty* list. This means the database hasn't been *modified* by the example.

## Step 04: Create and Verify a Keyspace

Inside the CQLSH shell, create a keyspace named `docker_hello_world`.

```
create keyspace docker_hello_world  
with replication = {  
    'class' : 'SimpleStrategy',  
    'replication_factor': 1  
};
```

### Explanation:

- **create keyspace:** Creates a new keyspace.
- **docker\_hello\_world:** The name of the keyspace.
- **with replication = { 'class' : 'SimpleStrategy', 'replication\_factor': 1 }:** Specifies the replication strategy and factor.

Verify the keyspace creation:

```
select *  
from system.schema_keyspaces  
where keyspace_name = 'docker_hello_world';
```

If the keyspace was created successfully, you will see the entry in the query result.

```
cqlsh> create keyspace docker_hello_world  
... with replication = {  
...     'class' : 'SimpleStrategy',  
...     'replication_factor': 1  
... };  
cqlsh> select *  
... from system.schema_keyspaces  
... where keyspace_name = 'docker_hello_world';  
  
keyspace_name | durable_writes | strategy_class | strategy_options  
-----  
docker_hello_world | True | org.apache.cassandra.locator.SimpleStrategy | {"replication_factor": "1"}  
  
(1 rows)  
cqlsh> █
```

## Step 5: Stop and Remove the Cassandra Container

Exit the CQLSH shell and remove the Cassandra container.

Quit

```
docker stop cass1
```

```
docker rm -vf cass1
```

```
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker stop cass1
cass1
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker rm -vf cass1
cass1
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$
```

### Explanation:

- **quit:** Exits the CQLSH shell.
- **docker stop cass1:** Stops the cass1 container.
- **docker rm -vf cass1:** Removes the cass1 container forcefully and deletes associated resources.



## Step 6: Test Data Recovery

Create a new Cassandra container and attach the volume to it.

```
docker run -d \  
  --volume cass-shared:/var/lib/cassandra/data \  
  --name cass2 \  
  cassandra:2.2
```

```
term@ubuntu-c1s35h-5cfccb9d45-1kb2w:~$ docker run -d \  
  --volume cass-shared:/var/lib/cassandra/data \  
  --name cass2 \  
  cassandra:2.2  
fbc1708dbd1e89bccfe6f37f2f420106054f650a83fee9fa0739b10e7ecda2f0  
term@ubuntu-c1s35h-5cfccb9d45-1kb2w:~$ docker ps  
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS                               NAMES  
fbc1708dbd1e   cassandra:2.2  "docker-entrypoint.s..."  53 seconds ago  Up 49 seconds  7000-7001/tcp, 7199/tcp, 9042/tcp, 9160/tcp  cass2  
term@ubuntu-c1s35h-5cfccb9d45-1kb2w:~$
```

**Connect to the new Cassandra container using CQLSH.**

```
docker run -it --rm \  
  --link cass2:cass \  
  cassandra:2.2 \  
  cqlsh cass
```

**Query the keyspace to verify data persistence.**

```
select *  
  
from system.schema_keyspaces  
  
where keyspace_name = 'docker_hello_world';
```

If the keyspace `docker_hello_world` is listed in the result, it confirms that the data persisted in the `cass-shared` volume

```

term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker run -it --rm \
--link cass2:cass \
cassandra:2.2 \
cqlsh cass
Connected to Test Cluster at cass:9042.
[cqlsh 5.0.1 | Cassandra 2.2.19 | CQL spec 3.3.1 | Native protocol v4]
Use HELP for help.
cqlsh> select *
... from system.schema_keyspaces
... where keyspace_name = 'docker_hello_world';

keyspace_name | durable_writes | strategy_class | strategy_options
-----
docker_hello_world | True | org.apache.cassandra.locator.SimpleStrategy | {"replication_factor":"1"}

(1 rows)
cqlsh>

```

## Step 7: Clean Up

Exit the CQLSH shell and remove the containers and volume.

`quit`

`docker rm -vf cass2`

`docker volume rm cass-shared`

## Explanation:

- **quit:** Exits the CQLSH shell.
- **docker rm -vf cass2:** Removes the cass2 container forcefully.
- **docker volume rm cass-shared:** Deletes the cass-shared volume.

```

term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker ps
CONTAINER ID   IMAGE      COMMAND                  CREATED        STATUS        PORTS                               NAMES
fbc1708dbd1e   cassandra:2.2  "docker-entrypoint.s..."  4 minutes ago  Up 4 minutes  7000-7001/tcp, 7199/tcp, 9042/tcp, 9160/tcp  cass2
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker volume ls
DRIVER    VOLUME NAME
local     cass-shared
local     e133c86aa234ae41ca0e53f58c281f8d84128698f370414f2033bd0e04e75b26
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker rm -vf cass2
cass2
term@ubuntu-c1sj5h-5cfccb9d45-1kb2w:~$ docker volume rm cass-shared
cass-shared

```

## Summary

### In this scenario, we have:

- Created a Docker volume.
- Ran a Cassandra container with the volume mounted.
- Connected to the Cassandra container using CQLSH.
- Created and verified a keyspace.
- Stopped and removed the Cassandra container.
- Tested data recovery by creating a new container and verifying the keyspace.
- Cleaned up by removing the containers and volume.

This demonstrates how to use Docker volumes for data persistence in a Cassandra database, ensuring that data remains available even after the container is deleted.